

How Uncertainty About Heterogeneity Impacts Technology Adoption*

[\[Click here for most recent version\]](#)

Hossein Alidaee[†]

12/8/22

Abstract

Individuals can learn about new technologies through peers or through more official sources. Peers' information is often based on only a handful of experiences. By contrast, official sources, such as the government, back their information with rigorous testing. In my setting of agricultural technology adoption, government recommendations are documented to be no more influential at inducing adoption than those of peers. This implies that data from peers is more influential per datum. I propose that this arises because returns to technology adoption are heterogeneous based on context and individuals face uncertainty about the context where government testing took place. I confirm this mechanism using a lab-in-the-field experiment with 1,571 farmers in Odisha, India. I also demonstrate that both survey data and results from a broad set of recent field experiments on agricultural extension are consistent with my mechanism that farmers place greater value on information with less context uncertainty.

*Thanks to Chris Udry, Ben Golub, and Lori Beaman for their guidance and support. I also thank Anna Waggener, Jaya Wen, Alireza Tahbaz-Salehi, Sneha Subramanian, Amanda Starc, Laura Schechter, Mark Rosenzweig, Bryony Reich, Daniel Rappoport, Ameet Morjaria, Daniel Martin, Paula Lopez-Peña, Pete Klibanoff, Erin Kelley, Dean Karlan, Seema Jayachandran, Olivia Healy, Xavier Gine, Nilesch Fernando, Erika Deserrano, Eddie Dekel, Arun Chandrasekhar, Ivan Canay, Meghan Busse, Federico Bugni, Zack Barnett-Howell, Bruno Barsanetti, Yuehao Bai, and Eric Auerbach for helpful conversations and comments, along with participants at the NU Development Lunch and Kellogg Strategy Brown Bag. This project was made possible by the research assistance of Uddesh Khuntia and the JPAL South Asia - Odisha field team, game development by Alejandro Quan-Madrid, and patient qualitative piloting with countless friends during the COVID-19 epidemic. I gratefully acknowledge financial support from Northwestern's Kellogg School of Management, Global Poverty Research Lab, and Buffett Institute for Global Affairs. This study is registered in the AEA RCT Registry under identifier AEARCTR-0009537. All errors are my own.

[†]Northwestern University. Email: hossein.alidaee@kellogg.northwestern.edu

1 Introduction

Adopting a new technology requires first learning about it: Is it effective? Is it reliable? Is it accessible? Because the technology is new, the answers to these questions are necessarily mired in uncertainty. One important source for this information is peers (Rogers 2003). Learning from our peers, or social learning, is an important information channel for decisions across domains, such as health (Qiao et al. 2020) or personal finance (Duflo and Saez 2002). Peers often have limited experience with a new technology: a friend may only experiment with the technology once or twice before sharing their experience. Despite this limited experience with a new technology, these peer recommendations are nevertheless an influential source of information.

Understanding why social learning is influential, despite peers’ limited information, is important for governments and NGOs designing their own campaigns to disseminate information. In the agricultural sector, agricultural ministries spend hundreds of millions of dollars on agricultural extension services per year (Sajesh and Suresh 2016), which introduce farmers to new technologies that have been vetted by extensive research. Yet, extension services are not necessarily more influential than social learning (Krishnan and Patnam 2014; Takahashi, Mano, and Otsuka 2019), wherein recommendations are fueled by very small scale experimentation by peers. Without an explanation for this disparity, social learning appears unreasonably influential.

This paper proposes one potential mechanism for why social learning is influential, which I refer to as *context uncertainty*. Context uncertainty arises from an agent’s inability to predict his own expected outcome from another individual’s recommendation, due to a lack of information about how their settings differ. This mechanism is rooted in two observations: (i) technologies have heterogeneous returns and (ii) individuals understand the characteristics of their peers better than those of other information providers. For example, I may know the difference in soil content between myself and my neighbor’s plot with a high level of certainty. Though my neighbor’s experience is limited, his recommendations carry great weight because I know how to extrapolate them to my own context.¹ By contrast, the recommendation I receive from an extension agent is based on testing in a distant location,

¹This paper presents the model under the assumption that the mapping required for extrapolation is known. For example, even though a farmer does not know how much less rainfall he has than a government test plot, he does understand how rainfall impacts the need for alternative irrigation technology. This is done for ease of exposition. However, the model extends to the case where an agent faces uncertainty about the mapping process as well.

on plot conditions I am unfamiliar with. Alternatively, the extension agent may not even know the test conditions himself, and is unable to share that information with the farmer. As a result, though the extension agent’s recommendation is generated by a bigger sample size, it suffers from high context uncertainty.

I test this mechanism using an experiment with a sample of 1,571 farmers in Odisha, India. The experiment involves each participant playing a mobile game about agricultural technology adoption. In each round of the game, the participant must decide how intensely he will adopt the hypothetical agricultural technology presented to him. Before making this decision, the participant has access to information from several fictional characters in the game, who experimented with the technology in the prior season and privately share their experiences. Each player plays multiple rounds of this game, with each round providing different information from the characters. This lab-in-the-field design allows me to isolate and vary context uncertainty as a parameter while holding all confounding factors, such as other forms of uncertainty, constant.

In the game, I tell participants that the fictional characters are one of two types: orange or blue. I also tell each participant that he himself is the orange type. Types represent a very simplified notion of context. Accordingly, all orange types have the same underlying returns to the technology. However, they may provide the respondent with different signals. This is due to idiosyncratic risk, i.e. poor implementation. This idiosyncratic risk is the source for sampling error. Analogously, blue types also share an underlying return to the technology. While the blue type has the same level of sampling error as the orange type, the technology is less suited to their needs, so they report lower returns to adoption. The participant is told exactly how much less effective the technology is for blue types, and is given a visual aid on how to adjust a signal from a blue type character to an equivalent signal from an orange type character.

Context uncertainty in the game arises from the fact a percentage of the fictional characters providing signals do not reveal their type. Instead, they appear as gray. Even though they are either orange or blue types, the participant does not know which. Consequently, the participant does not know whether the signal from these gray characters does or does not need to be adjusted for. In low context uncertainty rounds of the game, one quarter of characters will be gray. In high context uncertainty rounds, three quarters of characters will be gray. In each round, the empirical distribution of signals remains unchanged. This serves as my experimental variation for context uncertainty.

Though the game is detailed from the perspective of a designer trying to control for confounding variables, it contains very few rules, which are communicated clearly to the participant by local staff. Additionally, the participant is read a script that provides several examples to help motivate the different elements of the game. The game also begins with a training module of five rounds with distinct information environments. Data from this module is not recorded, and instead served to familiarize participants with the game and allow them several rounds during which they could ask questions. Heterogeneity analysis of the experimental data shows that context uncertainty aversion is unaffected by level of education, except at the lowest levels.

The experimental results are consistent with my theory: participants reduce their adoption intensity in rounds with higher context uncertainty. These results are robust to a non-parametric analysis of the experimental data. Further, I show that there is little heterogeneity in context uncertainty aversion based on demographic data collected at listing. These results suggest that context uncertainty aversion is a general phenomenon that is not driven by a specific subset of the population.

In addition to the experiment, each participant in my sample also participates in a survey about their prior agricultural technology adoption decisions. This survey data is consistent with my experimental results. For example, many farmers cite reasons pertaining to heterogeneity, such as soil differences, for why they have not put more weight on recommendations from their extension agent in the past. By contrast, less than five percent of farmers report corruption as one of their concerns with extension recommendations. I also find that citing heterogeneity in the past is correlated with higher levels of context uncertainty aversion within my experiment.

To further emphasize the external validity of context uncertainty aversion as a phenomenon, I also review several interventions previously studied in the literature on agricultural extension design. I show that context uncertainty is consistent with the results of a wide range of prior field experiments. I also explain how it may also be a factor in previously documented mechanisms in the technology adoption literature, such as the network centrality of initial adopters or homophily.

These results have implications for information provision more broadly, beyond the setting of agricultural technology adoption. Evidence of context uncertainty aversion suggests that designers of information campaigns should (i) provide information on heterogeneous returns when introducing technologies, and (ii) allocate more resources to acquiring data on hetero-

geneous returns instead of investing in more data from a narrow set of contexts. Notably, these two recommendations are complementary. Even if recommendations are personalized, an individual may still suffer from context uncertainty aversion if he either is not aware of the personalization or is concerned that the recommendation is missing information about his context.

To my knowledge, I am the first to provide a formal model of this mechanism and the first to provide well identified experimental evidence of its existence. Prior research on technology adoption informally discuss the role of individuals having concerns about the external validity of recommendations. Examples of this exist in agricultural technology adoption ([Munshi 2004](#)), climate policy ([Dechezleprêtre et al. 2022](#)), and medicine ([Alsan et al. 2022](#)). However, these papers provide neither a formal framework for exploring the implications of this mechanism nor are they able to isolate the mechanism from other factors.

The implications of this mechanism likewise contribute to a broad literature on the economics of information provision, which Haaland, Roth, and Wohlfart ([Forthcoming](#)) provides an extensive overview of. I show that whenever aggregated information is provided, agents have reason to discount the recommendation if they believe themselves to have characteristics varying from the mean of the distribution. While some of these interventions are certainly effective, a counter-factual intervention with personalized data could have greater impact. It is also important to provide contextual data to those who would consequently reduce adoption. If a source provides an individual with information suggesting better results than what the individual subsequently experiences, this can reduce an individual’s trust in future information from that source.

Despite this broad applicability, my experimental design, described in [Section 5](#), focuses on the agricultural technology adoption decisions of small and marginal farmers in rural India. The frictions behind agricultural technology adoption in developing countries is a longstanding focus of the development literature Foster and Rosenzweig ([2010](#)). This project is closely related to multiple strains of this literature.

First, I expand on several prior papers which investigate the role of heterogeneity on farmer productivity. For example, Suri ([2011](#)) shows that cross-sectional differences in technology adoption by Kenyan farmers can be explained by the technology’s heterogeneous returns. In related work, Gollin and Udry ([2021](#)) show that unobserved heterogeneity together with measurement error explain a large fraction of dispersion in agricultural productivity in Tanzania and Uganda. I demonstrate that heterogeneity not only impacts agricultural productivity

and technology adoption directly, but also through the impact that heterogeneity has on the uncertainty faced by farmers.

I also contribute to a second line of research that does study the impact of heterogeneity on farmer risk, but focuses on the adoption of micro-insurance services. This literature uses the term *basis risk* to describe the difference between the actuarial probability distribution used to price an insurance product and the actual risk distribution faced by a farmer. For example, for rainfall insurance, a rain gauge may be placed in the center of a district but conditions such as distance or heterogeneous terrain cause farmers to experience a different amount of rainfall on their own plots. Giné, Townsend, and Vickery (2008) document that insurance take-up is correlated with a decrease in rainfall basis risk and show that this is consistent with survey evidence among non-purchasers. This result is extensively corroborated by other work theoretical and empirical work on micro-insurance, such as Karlan et al. (2014), Clarke (2016), and Mobarak and Rosenzweig (2012). From this lens, the contextual differences I highlight between test plots and farmers’ own plots can be considered a form of basis risk that is reducing the take-up of agricultural technologies more broadly, beyond just micro-insurance products.

Third, my paper is directly relevant to, and consistent with, the findings of multiple empirical papers on the role of information on agricultural technology adoption. Section 8 reviews a variety of experiments on agricultural extension design specifically. There, I demonstrate external validity by illustrating that context uncertainty explains the efficacy of a broad swath of designs, in contrast to several alternative mechanisms.

In order to provide insight on how heterogeneity impacts uncertainty and learning, I provide a formal model for context uncertainty in Section 3. A burgeoning theory literature studies economic agents behaving under a variety of learning procedures. Examples include Liang (2020), Montiel Olea et al. (2022), and Salant and Cherry (2020). My model of learning is based on the standard Gaussian learning model, which has been used as a foundation for a variety of theoretical papers. This includes work on heterogeneous learning Haltiwanger and Waldman (1985). However, this work focuses on heterogeneity in priors or learning rules, rather than in the parameter of interest. A notable exception is Manski (2004). He studies heterogeneous preferences by agents who received imperfect data from peers in the sense of not observing counter-factual outcomes. While my mechanism can be generalized to the case of heterogeneous preferences, I assume that agents conduct small experiments, so uncertainty only arises due to sampling error and context uncertainty.

This model also contributes to a literature concerned with social learning specifically. The model presented here most closely resembles those in Sethi and Yildiz (2016) and Dasaratha, Golub, and Hak (2021). However, these papers focus on aggregate behaviors of the network, rather than individual decision makers. Sethi and Yildiz (2016) studies the setting closest to my own, but seeks to understand endogenous network formation when perspectives, rather than objective characteristics, are heterogeneous.

My work is also connected to the longstanding literature on portfolio theory. Though other work has considered agricultural technology adoption as a portfolio problem, I am the first to combine this with the context uncertainty mechanism. Within this literature, Gollier (2011) similarly studies the portfolio problem under multiple possible probability distributions. However, he uses the smooth ambiguity framework of Klibanoff, Marinacci, and Mukerji (2005) to study the case of decision making under ambiguity aversion. Although my results can be analogously obtained under the case of context uncertainty causing ambiguity in the smooth ambiguity model, my model makes no use of ambiguity. Instead, I model context uncertainty as one of multiple sources of risk.

Overview

The paper is organized as follows. Section 2 provides a brief intuition for the notion of context uncertainty. Section 3 provides a formal yet stylized model of learning in the presence of context uncertainty. Section 4 shows the implications of this model for technology adoption. Section 5 provides a detailed description of the lab-in-the-field experimental design. The results of the experiment are evaluated in Section 6. Section 7 supports the external validity of these findings using survey evidence from the same 1,571 farmers participating in the experiment. Section 8 shows that these findings are consistent with a variety of empirical results from field settings. Section 10 concludes.

2 An Intuition for Context Uncertainty

Context uncertainty is uncertainty arising from an agent’s inability to predict his expected outcome based on information from someone else, due to a lack of information about how their settings differ. The intuition can be illustrated by comparing three cases: (i) a world with homogenous returns to a technology, (ii) a world with heterogeneous returns and perfect information, and (iii) a world with heterogeneous returns and context uncertainty. In each case, I consider an agent i learning from two signals: (i) a peer j with limited experience

using the technology and (ii) an official recommendation from an extension agent k , derived from extensive experimentation.




Figure 1: Homogenous Risk

First, consider a world with homogenous types. Here, all agents are interested in learning a single value θ . Uncertainty arises only due to signal error: some individuals receive noisier signals about θ than others, as illustrated in Figure 1. Agent i 's peer j has only experimented with the technology on a single plot for a single season, so their signal about θ is noisy. As a result, j 's distribution in Figure 1a has a large variance. By contrast, the extension agent's signal in Figure 1b has a narrower variance, generated by more extensive testing. In this setting with homogenous types, agent i finds that the extension agent signal is more precise.

However, types may be heterogeneous. For an agricultural technology, types can represent differences in attributes like soil content, terrain, or opportunity costs for factors of production. When types are heterogeneous, returns are as well. Instead of learning about a single value θ , each agent is trying to learn about their individual return θ_j . This individual return is a mixture of the average return, θ , and an idiosyncratic component, γ_j , due to agent j 's type or *context*. Agent j 's return to adopting the technology is therefore

$$\underbrace{\theta_j}_{\text{Agent } j\text{'s Return}} = \underbrace{\theta}_{\text{Average Return}} + \underbrace{\gamma_j}_{\text{Agent } j\text{'s Context}}. \quad (1)$$

θ_j is not directly observed. Instead, agent j receives a noisy signal $s_j = \theta_j + \epsilon_j$ of this return, where the variance of the noise $\epsilon_j \sim \mathcal{N}(0, \sigma_j^2)$ is a function of experience.

When agent j shares s_j with an agent i , what does agent i learn about his own θ_i ? Agent i knows his own return is

$$\underbrace{\theta_i}_{\text{Agent } i\text{'s Return}} = \underbrace{\theta_j}_{\text{Agent } j\text{'s Return}} - \underbrace{\gamma_j}_{\text{Agent } j\text{'s Context}} + \underbrace{\gamma_i}_{\text{Agent } i\text{'s Context}}. \quad (2)$$

If agent i knows the difference $\gamma_i - \gamma_j$ between j 's context and his own, he can adapt the noisy signal s_j to $s_j^A = s_j - \gamma_j + \gamma_i$ to learn about his own returns. This adapted signal yields the same perception of risk as in the case of homogenous types. I illustrate this in Figure 2. Agent i can take the signal from peer j about θ_j , as in Figure 2a, and translate it to the signal in Figure 2c about θ_i . If agent i also knows how his context differs from the extension agent's, he can repeat this adaptation for his signal s_E , illustrated in figures 2b and 2d. Once the signals are translated, the case is analogous to the homogenous risk scenario of Figure 1. Agent i again finds the extension agent's signal to be more precise.

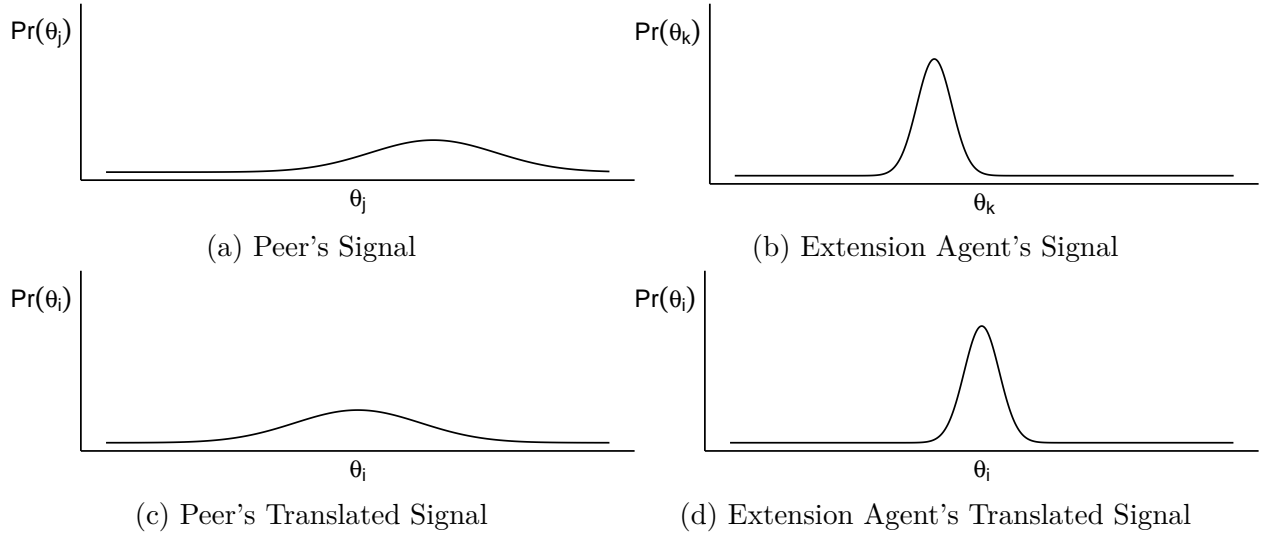


Figure 2: Heterogeneous Risk without Context Uncertainty

Context uncertainty arises when agent i does not know γ_j . Without this information, agent i cannot complete the translation in Equation 2. He may, however, know a distribution of possible values for γ_j . The variance in possible values for γ_j is the *context uncertainty*. The amount of context uncertainty is a key difference between information from peers and information from extension agents. Figure 3 provides a stylized example of this. Agent i knows his peer agent j extremely well, and therefore knows γ_j exactly. Consequently, figures 3a and 3c are identical to figures 2a and 2c.

By contrast, agent i does not know the extension agent's context γ_k precisely; he is not given information about the weather patterns or soil content of test plots. However, agent

i may know several possible values for γ_k and the probability of each. For example, he knows that the village on the other side of the mountain is less prone to flooding. This imperfect information about γ_k means that the extension agent’s signal about θ_k , illustrated in Figure 3b, is translated into a distribution of possible distributions about θ_i . This is illustrated in Figure 3d; there are multiple distinct possible distributions for θ_i . Agent i has some belief about the probability of each distribution being true. That probability corresponds to his belief about the value of γ_k .

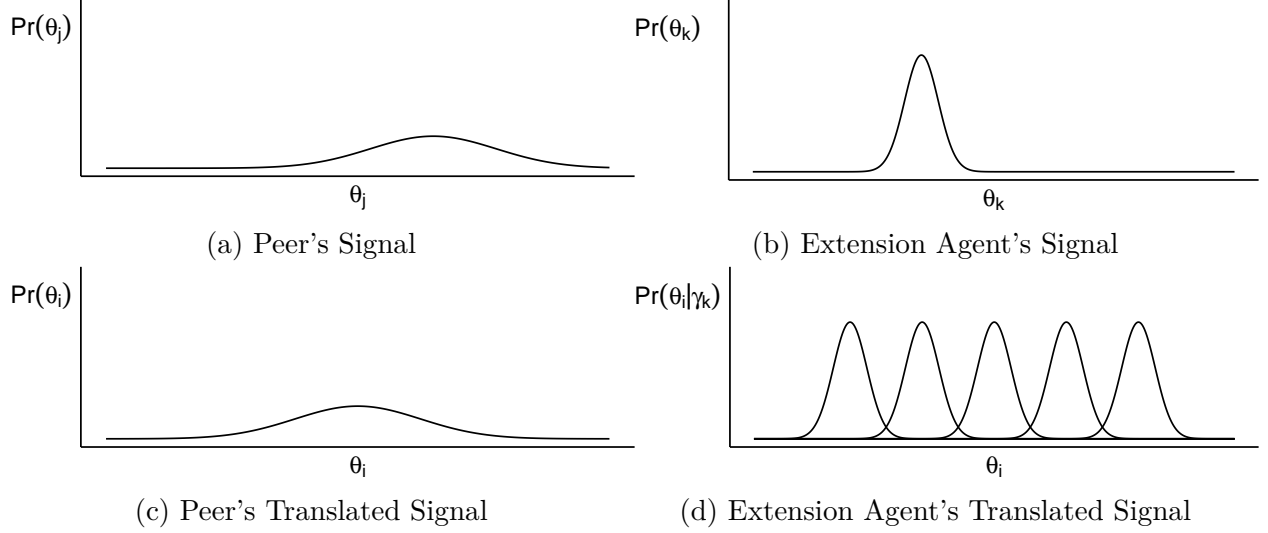


Figure 3: Heterogeneous Risk with Context Uncertainty

Heterogeneity causes context uncertainty when an agent does not know enough about the differences to translate signals from another context to his own. Figures 3c and 3d show that this additional uncertainty means that agent i no longer perceives the extension agent’s signal to be more precise than peer j , whose context is well known.

The distinctive distribution in Figure 3d arises from agent i ’s discrete beliefs about γ_k , and chosen to convey the impact of heterogeneity. In Section 3, agent i ’s beliefs about γ_k form a normal distribution, amounting to an extra Gaussian noise term. Consequently, the extension agent’s translated signal will also be Gaussian, but with a large variance parameter. Figure 4 illustrates one example of this.




Figure 4: Translated Signals With Gaussian Context Uncertainty

3 A Model of Learning from Heterogeneous Sources

An agent A is trying to learn about a new technology. He is interested in identifying his personal marginal return to adopting the technology, relative to the status quo technology already being commonly used. This marginal return is unobservable and denoted by $\theta_A \in \mathbb{R}$. As described in Equation 1, his marginal return is a function of the average marginal return $\theta \in \mathbb{R}$ and his context $\gamma_A \in \mathbb{R}$. Agent A knows the value of his own context γ_A . However, the average marginal return $\theta \in \mathbb{R}$ is also not observable.

3.1 The Information Environment

Agent A does not experiment with the new technology, so he receives no private signals about its efficacy. He does, however, receive private signals from a set of peers $N = \{1, \dots, n\}$. He also receives a signal from an extension agent E . To avoid focusing on issues of social learning patterns, we assume the peers do not share information with each other or the extension agent, and that the extension agent also only shares information with agent i . Equivalently, the village is a social network G whose topology is a directed star graph S_{n+1} , as illustrated in Figure 5.

Each individual $j \in \{1, \dots, n, E\}$ receives a single private signal s_j about their own marginal return to the technology, $\theta_j \in \mathbb{R}$. This noisy signal can be decomposed as

$$s_j = \theta + \gamma_j + \epsilon_j. \quad (3)$$

where $\epsilon_j \sim \mathcal{N}(0, \sigma_j)$ are independent normally distributed random variables. We refer to ϵ_j as agent j 's *sampling error* and σ_j as their *sampling uncertainty*. Further, we will refer to any independent Gaussian random variables with mean zero and finite variance, including




Figure 5: Signal Sharing Network S_{n+1}

ϵ_j , as *Gaussian white noise*. Comparing Equation 1 with Equation 3 clarifies that a signal is a noisy observation of θ_j , where noise arises from limited experience testing the technology.

Each of these peers and the extension agent subsequently shares their signal with agent A . They also share their values σ_j by communicating the sample size behind their recommendation.² However, agent A does not know the context for other individuals equally well. His belief about each context is normally distributed

$$\gamma_j \sim_A \mathcal{N}(\mu_j^\gamma, \sigma_j^\gamma)$$

where mean μ_j is the *expected context* and σ_j^γ is the *context uncertainty* A has about j . Because the focus of this paper is on context uncertainty parameter σ_j^γ , we assume that agent A 's estimate of μ_j is known to be unbiased so that $\mu_j^\gamma = \gamma_j$ and is believed with certainty.³⁴

²This can be generalized to also sharing other attributes behind sampling error, such as the care placed into their experimentation.

³A possible extension of this model is introducing additional uncertainty from agent A 's estimation μ_j^γ .

⁴An alternative specification can assume that the context for each individual is drawn from a personal distribution $\mathcal{N}(0, \sigma_j^\gamma)$. In this setting, we assume that agent A knows the distribution for each agent, but does not know their realization. This specification implies that agent A will not need to subtract a mean belief to evaluate his translated signal in Equation 4 of Section 3.2. Instead, he can allow $s_j^A = s_j + \gamma_i$, where his own context γ_i is a known scalar and his beliefs over γ_j act as another centered Gaussian noise parameter, as with ϵ_j . This amounts to skipping directly to Equation 5.

3.2 Signal Translation

In order to use the information from his peers, agent A first must translate their signals to his context in order to learn about his own marginal return, θ_A . Recall that, for each agent, he observes their signal s_j and sampling error σ_j and also has beliefs over their context parameterized by μ_j^γ and σ_j^γ . He uses this information to create a modified signal

$$s_j^A = \underbrace{\theta + \epsilon_j + \gamma_j}_{\text{Original Signal}} + \underbrace{(\gamma_i - \mu_j)}_{\text{Context Adjustment}} \quad (4)$$

composed of the original signal he received from an individual j and adjusting it using his own γ_i and his belief μ_j . This amounts to agent A receiving a signal with the structure

$$s_j^A = \theta + \epsilon_j + \bar{\gamma}_j + \gamma_i \quad (5)$$

where $\bar{\gamma}_j \sim_A \mathcal{N}(0, \sigma_j^\gamma)$ is a random variable denoting agent A 's uncertainty over the difference between agent j 's realized context and his own γ_i . Under this translation, agent A interprets $\bar{\gamma}_j$ as a Gaussian white noise parameter.

Note that equations 4 and 5 imply that s_j^A is the sum of two scalars, θ and γ_i , and two Gaussian random variables, $\epsilon_j \sim \mathcal{N}(0, \sigma_j^2)$ and $\bar{\gamma}_j \sim_A \mathcal{N}(0, (\sigma_j^\gamma)^2)$. The arithmetic for Gaussian random variables implies that

$$s_j^A \sim_A \mathcal{N}(\theta + \gamma_i, \sigma_j^2 + (\sigma_j^\gamma)^2). \quad (6)$$

The Gaussian structure of each translated signal s_j^A , with common mean $\theta + \gamma_i$, allows agent A to cleanly aggregate signals from various sources in the next subsection. A similar analysis of a system of Bayesian signals can be done with any conjugate distribution.⁵

3.3 Learning About his Own Return

Once translated, Agent A can use these modified signals, s_j^A for $j \in \{1, \dots, n, E\}$ to form his belief about θ_A . By Bayes' Rule, agent A 's posterior belief about θ_A is proportional to his

⁵See Rossi, Allenby, and McCulloch (2012) for an overview of conjugate distributions, solving these models analytically, cases when deference to numerical integration methods are needed, and how to do so.

prior belief times the conditional likelihood of receiving his signals:

$$\Pr(\theta_A | s_1^A, \dots, s_n^A, s_E^A) \propto \Pr(s_1^A, \dots, s_n^A, s_E^A | \theta_A) \Pr(\theta_A).$$

Because he has little knowledge about the technology, we assume that agent A 's prior for his marginal return to adoption θ_A is

$$\theta_A \sim_A \mathcal{N}(0, \sigma_0^2)$$

where σ_0 is sufficiently large so that the prior is weak and minimally informative.⁶

It follows from Bayes rule for linear Gaussian systems⁷ that agent A 's posterior is distributed $\mathcal{N}(\tilde{\mu}, \tilde{\sigma}_0^2)$ where

$$\tilde{\mu} = \tilde{\sigma}_0^2 \left(\frac{0}{\tilde{\sigma}_0^2} + \frac{s_E^A}{\sigma_E^2 + (\sigma_E^\gamma)^2} + \sum_{j \in 1, \dots, n} \frac{s_j^A}{\sigma_j^2 + (\sigma_j^\gamma)^2} \right) \quad (7)$$

and

$$\tilde{\sigma}_0^2 = \left(\frac{1}{\tilde{\sigma}_0^2} + \frac{1}{\sigma_E^2 + (\sigma_E^\gamma)^2} + \sum_{j \in 1, \dots, n} \frac{1}{\sigma_j^2 + (\sigma_j^\gamma)^2} \right)^{-1}. \quad (8)$$

We note that the posterior mean is the weighted average of the prior mean, the extension agent's translated signal, and the translated signal from each peer. An important observation is that the weight on each term is determined not only by the noise arising from sampling error, σ_j , but also the noise from context uncertainty, σ_j^γ . These context uncertainty values σ_j^γ also contribute to increasing the overall variance $\tilde{\sigma}_0^2$ of the posterior belief.

⁶Alternatively, we can consider agent A 's prior to be the improper prior, derived from the limit

$$\lim_{\sigma_0 \rightarrow \infty} \mathcal{N}(0, \sigma_0^2).$$

Because the maximum likelihood estimator is the maximum a posteriori estimator (Murphy 2012), agent A 's posterior will be equivalent to the computed likelihood distribution under the limit. The parameters in equations 7 and 8 can then be replaced with their limits:

$$\tilde{\mu} = \tilde{\sigma}_0^2 \left(\frac{s_E^A}{\sigma_E^2 + (\sigma_E^\gamma)^2} + \sum_{j \in 1, \dots, n} \frac{s_j^A}{\sigma_j^2 + (\sigma_j^\gamma)^2} \right) \quad \text{and} \quad \tilde{\sigma}_0^2 = \left(\frac{1}{\sigma_E^2 + (\sigma_E^\gamma)^2} + \sum_{j \in 1, \dots, n} \frac{1}{\sigma_j^2 + (\sigma_j^\gamma)^2} \right)^{-1}.$$

⁷See Section 4.4 of Murphy (2012) for a more complete overview of linear Gaussian systems and their properties.

These posterior parameters contain this paper’s primary message. It is evident, once we consider context uncertainty, why individuals do not learn more from sources such as extension agents. It is true that rigorous testing behind an extension agent’s recommendation constitutes lower sampling uncertainty σ_E^2 . However, his signal’s total uncertainty $\sigma_E^2 + (\sigma_E^\gamma)^2$ may exceed those of peer villagers. Peer villagers may know one another’s contexts extremely well, bringing context uncertainty as low as $\sigma_j^\gamma = 0$.

Given this, which source would a farmer listen to: his set of peers or the extension agent? Proposition 3.1 explains that, analogous to the issue of weighting signals in his posterior, the agent will always choose the signal with less total variance when forced to choose.

Proposition 3.1. *Let an agent A be a Bayesian expected utility maximizer choosing between the two signals:*

$$s_j^A \sim_A \mathcal{N}(\theta + \gamma_i, \sigma_j^2 + (\sigma_j^\gamma)^2) \quad \text{and} \quad s_E^A \sim_A \mathcal{N}(\theta + \gamma_i, \sigma_E^2 + (\sigma_E^\gamma)^2).$$

The agent will choose the signal with lower total variance.

However, Proposition 3.1 does not imply that a higher level of adoption will be chosen. The following section will explain the adoption problem and provide the comparative statics of adoption after imposing the additional assumptions needed on our agent.

4 How Context Uncertainty Impacts Technology Adoption

Context uncertainty’s impact on learning can be interesting in isolation. However, this paper’s ultimate goal is understanding its ultimate impact on technology adoption behavior. Further, measuring respondents’ probabilistic beliefs is an active area of research, in both economics (e.g. Manski 2004; Enke and Graeber 2019) and human-computer interaction (e.g. Koval and Jansen 2022; Greis et al. 2017; Greis et al. 2019), concerned with measurement error issues that could obfuscate real differences between priors and posteriors. Consequently, this section explains the impact of learning on the agent’s technology adoption problem and derives the main results. Section 6 tests these propositions about adoption levels empirically.

4.1 The Agent's Problem

How does agent A use the signals shared with his by his peers and extension agent? He ultimately seeks to maximize his plot's yield by choosing the right level of adoption of the new technology, based on his beliefs about θ_A .

The agent's yield $Y_A(\alpha, \theta_A)$ is a function of the marginal impact θ_A of the new technology and the share of his plot $\alpha \in [0, 1]$ on which he adopts it. Its functional form is the sum of his yield on the share of land where the technology was adopted, and is θ_A units higher (lower), and his yield on the land using the status quo technology:

$$Y_A(\alpha, \theta_A) \equiv \alpha(1 + \theta_A + \nu_A + \epsilon_{A,1}) + (1 - \alpha)(1 + \nu_A + \epsilon_{A,0}). \quad (9)$$

Yield is subject to two forms of shocks. $\nu_A \sim \mathcal{N}(0, \sigma_V)$ denotes a common shock to all parts of his land, such as unexpected weather conditions for the season. $\epsilon_{A,1} \sim \mathcal{N}(0, \sigma_A)$ and $\epsilon_{A,2} \sim \mathcal{N}(0, \sigma_A)$ are identically distributed shocks that are independent across the two fractions of land. These shocks can be interpreted as the average outcome of implementation idiosyncrasies, such as accidental variation in pesticide application.

We assume that the risk averse agent A has a utility function

$$u(Y_A(\alpha, \theta_A))$$

where u is nondecreasing and concave. His goal is to choose a level of adoption α^* that maximizes the his expected utility over his distribution of possible yields:

$$\alpha^* = \arg \max_{\alpha} E[u(Y_A(\alpha, \theta_A))]. \quad (10)$$

Agent A 's posterior beliefs over θ_A follow the distribution $\mathcal{N}(\tilde{\mu}, \tilde{\sigma}_0^2)$ with parameters as defined in equations 7 and 8. To focus on the impact of context uncertainty on adoption and ignore issues of particular signal draws, we assume that $s_j^A = \bar{\mu} \in \mathbb{R}^{++}$ for all $j \in \{1, \dots, n, E\}$ so all individuals have the same positive signal draw and the mean of the belief distribution is $\tilde{\mu} = \bar{\mu}$.⁸

⁸The true marginal return to the technology is perhaps the optimal candidate for $\bar{\mu}$. It coincides with the expected value of each translated signals s_j^A . This is not set explicitly, both for the purpose of generality and to avoid confusion in notation between the true value of θ_A and agent A 's beliefs about the parameter.

4.2 Adoption Dynamics

Our principal goal is showing that the agent’s optimal level of adoption is a function of beliefs over θ_A , particularly that α^* weakly increases when context uncertainty decreases for any signal s_j^A received. This result is formalized in Proposition 4.1.⁹ A stronger version of this result can be obtained for strictly risk averse agents. The proof for this case is analogous.

Proposition 4.1. *Consider a risk averse agent A with utility u , which is nondecreasing and concave, selecting a level of adoption to solve Equation 10. The agent’s optimal level of adoption α^* is nonincreasing in context uncertainty from any signal.*

So far, we have restricted our focus to the impact of context uncertainty in isolation. However, it is also important to understand the interaction between context uncertainty and uncertainty arising from sampling error. It turns out that, for certain agents, these two forms of uncertainty are decision complementary, as defined in Definition 4.1¹⁰: when the agent faces less of one form of uncertainty, the marginal reduction in the other form of uncertainty has greater impact on α .

Definition 4.1. Two inputs x and y are **decision complementary** if, for an agent selecting an optimal α to maximize their expected utility, the optimal decision rule α^* has increasing differences in x and y :

$$\alpha^*(x'', y'') - \alpha^*(x'', y') \geq \alpha^*(x', y'') - \alpha^*(x', y')$$

for all $x'' \geq x'$ and $y'' \geq y'$.

As Proposition 4.2 documents, decision complementarity requires stronger assumptions about our agents preferences. For differentiable utility, this assumption is a restriction on the third derivative, U''' . This captures the skewness of the utility function. However, utility

⁹Proposition 4.1 makes provides sufficient, but not necessary, conditions for this comparative statics result. However, these weaker conditions require greater exposition. Readers interested in weakening these conditions on the utility function, or on signal’s distribution, should consult Section 4 of Athey (2002).

¹⁰This definition of complementary is stronger than pairwise increasing differences in the utility function U amongst parameters α , σ_j , and σ_j^γ . The difference is most easily understood in the continuous case. Pairwise increasing differences is a statement about the set of second partial derivatives $\frac{\partial^2 U}{\partial \alpha \partial \sigma_j}$, $\frac{\partial^2 U}{\partial \alpha \partial \sigma_j^\gamma}$, and $\frac{\partial^2 U}{\partial \sigma_j \partial \sigma_j^\gamma}$. By contrast, Definition 4.1 is a statement about the third derivative $\frac{\partial^3 U}{\partial \alpha \partial \sigma_j \partial \sigma_j^\gamma}$.

need not be differentiable. Definition 4.2 provides a general definition for our restriction, allowing us to continue using the lattice toolkit of monotone comparative statics.

Definition 4.2. Let $g_\epsilon(x) = x - \epsilon$. An agent’s utility u satisfies **decreasing absolute risk aversion (DARA)** if the function $v(x) \equiv u(g_\epsilon(x))$ is a concave transformation of $u(x)$ for all $\epsilon > 0$.

Proposition 4.2. *Consider a risk averse agent A with DARA utility u , and is selecting a level of adoption to solve Equation 10. For any given signal s_j , context uncertainty $(\sigma_j^\gamma)^2$ and sampling uncertainty σ_j^2 are decision complementary.*

These theoretical results have important implications for policymakers encouraging adoption of technologies. First, if the target of an information campaign is risk averse, then the information campaign should provide contextual data about where the testing occurred. For example, provide details about the plot traits such as soil nutrition and weather relative to the plot of the targeted individual.

Second, for policymakers deciding whether their information campaigns will be more influential with greater data, the answer depends on the agent’s preferences. If individuals exhibit decreasing absolute risk aversion, then the marginal impact of additional data is greater when more information about context is revealed. Depending on the cost of providing additional context data from existing testing versus conducting more testing, an information campaign can mimic the effect of additional testing just by revealing context.

5 Experimental Design

This lab-in-the-field experiment recruited 1,571 farmers in Odissa, India. Testing context uncertainty as a mechanism requires introducing variation that does not affect other, related factors. I tackle this difficulty by designing a mobile game where I completely control the information available.

The game has a simple premise: the farmer must decide how intensively to adopt a hypothetical technology on his plot. Before making his decision, the farmer has access to information from several fictional characters. Each character tested the technology in the previous season

and privately gives the farmer their impression. The farmer’s real payout is based on how effectively his own yield is maximized by choosing the right level of technology adoption.

The experiment proceeded as follows. First, respondents participated in a survey to collect information about their perception of peers, extension agents, and other sources of information about farming. After the survey was completed, an enumerator provided each participant with a scripted explanation of the game and compensation structure. The participant then begins by playing a five round practice module of the game, during which they are allowed to ask any questions about the game. At the end of the practice module, all remaining rounds are played in the order randomly assigned to the participant. When the player has completed all rounds, they are paid based on their performance and their participation is complete. Results of the game are only revealed through the final compensation to avoid ongoing learning.

The following subsection details the game’s overall design. Subsequently, I describe the individual game modules used to create different forms of identifying variation. Then, I describe the randomization procedure. The final subsection provides characteristics of the participants in the sample.

5.1 Game Design

This subsection provides details on the game design. All the information in this subsection is also provided to the participant.

Setting

The participant plays a multiple-round game. In each round of the game, he lives in a hypothetical village with hypothetical fellow villagers, referred to as *characters*. He is considering adoption of a hypothetical agricultural technology. The technology, village, and the characters are different in each round. In each round, he has not yet adopted this technology, and has no direct experience with it. However, some of the other characters in the game experimented with the technology last season. The participant’s goal is to use the information shared with him by these characters to decide how intensively to adopt the new technology on his own plot.

Signals

Each character privately shares his experience with the technology with the participant. These private signals are shared using emojis from a 7-point Likert scale, reflecting their perceived value of the technology over the status quo alternative.¹¹ The scale is illustrated in Figure 6.



Figure 6: Signal Likert Scale

Character Types

There are two types of villagers: the ‘orange’ type and the ‘blue’ type. The participant is the orange type. Types denote contexts: two characters with the same type have all the opportunity costs, soil patterns, and other factors that could impact a technology’s efficacy. Accordingly, the returns to the technology are identical within type.

However, signals are not identical across types. The hypothetical technology has less value for blue types. This disparity is explained using an example about rainwater catchment technology. In the example, blue types represent villagers with better alternative water access (e.g. adjacency to a river), while orange types represent those on more arid land. This highlights that the underlying value of the technology is unchanged, but the returns are heterogeneous based on plot type.

The participant is told the exact magnitude of this difference varies per module. All modules used for our main hypotheses provide full information about the translation. Section 5.2 provides additional details about individual modules.

Sampling Error

Two characters with the same type may share different signals with the participant. These differences arise due to *sampling error*. Sampling error is explained as a consequence of

¹¹There is a recent, extensive literature in the field of human-computer interaction on how to best visualize data, including visualizing uncertainty. One recent overview of trade-offs in this literature can be found in Kim, Moritz, and Hullman (2021). This project optimized visualization for its settings, where participants were not ubiquitously familiar with bar charts and had wide heterogeneity in education. While past research does feature concerns about emoji use, these concerns impact cross-sectional analysis (i.e. Alismail and Zhang 2020; Lu et al. 2016). I avoid these concerns by using a within-subject design, wherein I study the impact on differences in each participant’s technology adoption across rounds.

issues such as implementation error. These differences would disappear with greater testing. However, each character has only a single season of experience with the technology on a single plot of land.

Orange types experience the same amount of sampling error as blue types. The signal distribution for one type is therefore a scalar shift of the signal distribution for the other type. The scalar shift matches the difference in returns. For example, if the difference in returns across types is two units, the empirical distribution of signals seen from orange characters will be identical to that of blue characters, but shifted up by two units.

Context Uncertainty

Some fraction of characters will share their signal but their type is unknown. Visually, these characters appear as a gray, instead of blue or orange. The participant still sees the signal from these gray characters, but does not know whether their type is blue or orange. This is illustrated in Figure 7.




Figure 7: A Gray Character Has Context Uncertainty

A gray character is equally likely to be a blue or orange character. Consequently, the blue and orange distributions remain equivalent under a scalar shift. The percentage of gray characters will vary across rounds. This variation, within users but across rounds, is the main source of identification for testing behavior towards context uncertainty.

Adoption Decision

The participant's plot features ten rows. He must decide, given the information presented in a round, how many rows to adopt the hypothetical technology on. To ease the mapping

between signals and adoption decisions, a Likert Scale is provided underneath the sliding tool used to select adoption intensity, as illustrated in Figure 8.



Figure 8: Sliding Scale For Selecting Adoption Intensity

Payout

The participant's total compensation for playing the game is between 100 and 300 rupees, depending on performance. Compensation is designed to average around 200 rupees, approximately a half day's wage. In each round, yield is maximized by selecting the intensity equal to the true underlying mean signal for the orange type. The mobile game computes the absolute deviation in each round and determines total compensation by adding performance across all rounds. The total compensation owed is provided once the participant completes the game.

5.2 Module Descriptions

The game is divided into seven modules. The first module consists of five rounds and is unincentivized; it serves as a practice module to familiarize participants with the game. All remaining modules are incentivized and consist of only two rounds.

I test my two primary hypotheses, the comparative statics results derived in Section 4, using only two of the seven modules. The remaining modules test various possible non-neoclassical behaviors outside the scope of my model, but that may affect participant behavior. These are described in Appendix Section B and analyzed in Appendix Section C.

Context Uncertainty Under High Signal Error Module (CU:HSE)

In this module, participants are told that the technology works exactly two units better for orange types than blue types. This is visualized in Figure 9. I implement high signal error by setting a wider distribution of possible emoji signals within each type.¹² Both rounds feature 40 characters, and an equal number of blue and orange type characters. One round features 30 characters with unknown (gray) type out of the 40 total characters. The other round features only 10 characters with unknown type. The two rounds are illustrated in Figure 10.



Figure 9: Signal Adjustment Across Types



(a) Low Context Uncertainty



(b) High Context Uncertainty

Figure 10: Rounds in the Context Uncertainty Under High Signal Error Module

Context Uncertainty Under Low Signal Error Module (CU:LSE)

This module is identical to the previous module, except that signal error is lower. The two rounds are illustrated in Figure 11. As discussed in Section 5.3, the order of the rounds within and across modules is not kept constant; it is randomized across participants.

¹²The maximum signal error is limited by only having seven possible signal values. Signals have to be translatable across types, and signals translate by two units across types. Consequently, the distribution of signals within each type is limited to existing on five units.



(a) Low Context Uncertainty



(b) High Context Uncertainty

Figure 11: Rounds in the Context Uncertainty Under Low Signal Error Module

5.3 Randomization

My experimental design relies on identifying differences in adoption across rounds. To ensure I measure the impact of my variation, and not confounding factors, my experimental design randomizes the order of (i) the rounds being played, (ii) the technologies being adopted, and even (iii) the names of the hypothetical villages. The order of each is randomized across participants and each is randomized separately.

Round order is perhaps the most important dimension for randomization. This is for multiple reasons. First, one could expect participants learning across rounds thus having different behavior in later rounds. Avoiding this requires randomizing how many rounds appear before the round of interest. Second, relative round order may also be important. For example, a participant may respond differently to a high context uncertainty round depending on whether or not it was preceded by the low context uncertainty round.

To avoid either of these channels, I randomize both the order of modules and the order of rounds within modules. Following the recommendation of McKenzie (2022), I implement a matched quartet design. The index assigned to each participant is based on data summarized in the following subsection on sample characteristics and collected during an initial listing survey. A detailed algorithm is provided in Section D.

5.4 Sample Characteristics

My experiment takes place in the state of Odissa, India. A listing survey was conducted in two districts, Khorda and Cuttack. During the listing, field staff identified 1,868 potential

respondents meeting a set of five criteria. First, all individuals had to own five or fewer acres of land to meet the criteria of our population of interest, small or marginal farmers. Second, the individual needed to be the primary decision maker for their plot of land, and thereby be in charge of technology adoption decisions. A third restriction was placed on age: respondents could be a maximum of 60 years old. Early piloting suggested that individuals above this age often deferred actual control to younger members of the household, such as a son. Fourth, I also excluded individuals who did not primarily cultivate their own land, as sharecroppers face different incentives when deciding whether to adopt a new agricultural technology. Finally, potential respondents needed to be in contact with their extension agent, locally known as a Village Agricultural Worker (VAW), so that they could answer questions about how they perceive the VAW compared to other information sources.

Of the 1,868 identified individuals, 27 were dropped due to incomplete listing data. Table 1 summarizes the mean and standard deviation of the remaining individuals across a variety of attributes collected during the listing survey. The remaining individuals were randomly assigned to one of 64 treatment arms according to Section 5.3. 25 individuals per arm (1,600 in total) were designated as *primary* respondents at random. The remaining respondents were designated as backup respondents in the event that the enumeration team could not reach a primary respondent.

Table 1 illustrates a non-trivial amount of heterogeneity within our sample. This heterogeneity was used in two ways. First, these variables are used in Section 5.3 to create an index for stratified randomization. Second, Section 6 explores the heterogeneous impact of context uncertainty along these attributes.

Table 1: Sample Characteristics

Characteristic	Mean (SD)
Age	44.48 (8.14)
Monthly Expenses (INR)	6,628.27 (4,052.09)
# Of Dependents	2.79 (1.03)
# Of Crops Grown In Last 12 Months	4.40 (3.43)

% Employed In Village During Lean Season	59.6
% Migrating Elsewhere During Lean Season	19.4
% Self-identifying As Early Adopters	50.0
% Married	93.5
% Completed School Until 12th Standard	24.7
Observations	1,841

Notes: This table shows means of key attributes for all eligible individuals from the listing survey. Standard deviations are in parenthesis. 1,600 individuals from this sample were designated as primary households. The remaining households were designated as backup households during randomization. Monthly expenses winsorized at 99th percentile.

6 Results

This section tests the reduced form implications derived from the model from Section 4 using data from the lab-in-the-field experiment. Because the participant population consists of small and marginal farmers, I assume that the minimal assumption imposed by the theoretical propositions are met. Namely, evidence shows that these farmers exhibit not only risk aversion but also decreasing absolute risk aversion (DARA) (Bar-Shira, Just, and Zilberman 1997).

Testing Proposition 4.1: Does Context Uncertainty Reduce Adoption?

The first hypothesis is that risk-averse individuals reduce adoption as context uncertainty increases. My results are consistent with this hypothesis: within each module, farmers adopt more intensively in rounds with lower context uncertainty.

Testing the hypothesis is straightforward under my experimental design. In each of the two modules discussed in Section 5.2, there were two maps: one with high context uncertainty and one with low context uncertainty. Exploiting this, we can use the following regression specification:

$$\text{Adoption Choice}_{i,t} = \beta_0 + \beta_1 \cdot \mathbb{1}(t = \text{High Context Uncertainty}) + u_i + \epsilon_{i,t}.$$

In this specification, $\mathbb{1}(t = \text{High Context Uncertainty})$ is a dummy variable for whether the data was collected in the high context uncertainty round. u_i is an individual level fixed

effect, which is observed because each participant plays each round.




Figure 12: Coefficients for effect of increasing context uncertainty on adoption intensity

The hypothesis here is directional: if context uncertainty reduces adoption, $\beta_1 < 0$. I test this hypothesis with data from three cases. First, I test it using data from the high signal error module. Second, I test it separately using data from the low signal error module. Third, I pool the data from both modules.¹³ The results of these regressions are illustrated in Figure 12. A table is also provided in Table A1. All three specifications show evidence that context uncertainty reduces adoption.

The results of these three specifications each reject the null hypothesis. Across all specifications, I find evidence that farmers reduce their adoption intensity by approximately 0.3 rows when comparing a high context uncertainty environment to an environment with low context uncertainty. In the low signal error module, the mean level of adoption under low context uncertainty was 5.217. This indicates that farmers reduced their adoption by approximately 6% when they instead face high context uncertainty in the same module.¹⁴ However, given

¹³Pooling across modules with different levels of signal error does not impact interpretation of magnitudes. This is because the high context uncertainty round in both modules are identical in all aspects except within-type variance of emoji, and the same is true for the low context uncertainty rounds. As a result, the change in the level of context uncertainty is the same across modules: participants see either 10 or 30 characters with unknown (gray) type.

¹⁴The results are empirically similar in the high signal error module. In that module, mean adoption under low context uncertainty was 5.137. It was reduced 0.35 rows in the high context uncertainty round,

this result is within a lab setting, emphasis should be placed on the consistent qualitative effects.

How did context uncertainty impact the distribution of adoption? Figure 13 provides histograms of adoption intensity by participants in each module, comparing adoption under high versus low context uncertainty. The results of the one-sided Kolmogorov–Smirnov test in each module are significant. Further, it is possible visually inspect and see the distribution shifting towards lower levels of adoption intensity under higher context uncertainty in each module.




Figure 13: Histograms of Adoption Intensity by Level of Context and Signal Uncertainty

These results cumulatively suggest that the main hypothesis holds: higher context uncertainty leads to lower levels of adoption.

Future work in this section will explore the heterogeneity of this effect. Using gathered data on traits such as level of consumption and education, I will explore the relationships between these traits and in-game behavior within these modules.

indicating a 7% reduction in adoption intensity.

Table 2: Difference in Difference Design for Testing Complementarity

Sampling Error	Context Uncertainty		
		Low	High
	Low	α_{LL}	α_{LH}
	High	α_{HL}	α_{HH}
	Difference	$\alpha_{HL} - \alpha_{LL}$	$\alpha_{HH} - \alpha_{LH}$
		Difference	
		$(\alpha_{HH} - \alpha_{HL}) - (\alpha_{HL} - \alpha_{LL})$	

Testing Proposition 4.2: Are Context Uncertainty and Sampling Error Complements?

The second hypothesis has to do with the relationship between context uncertainty and sampling error. A more general interpretation of Proposition 4.2 is that the relationship between context uncertainty and sampling error have on jointly impacting adoption intensity depends on an individual’s risk preferences. If the individual has decreasing absolute risk aversion (DARA), the theorem states that these two forms of uncertainty are complements. Alternatively, under increasing absolute risk aversion (IARA), these two forms of uncertainty are substitutes. Finally, if an individual’s absolute risk aversion is constant (CARA), we expect the two forms of uncertainty to be independent.

Following prior work on testing complementarities [e.g. Athey and Stern (1998); Cassiman and Veugelers (2006)]¹⁵, I test for complementarity using a difference-in-difference approach. This approach is illustrated in Table 2, where each α parameter denotes the level of adoption chosen by participants in a given information environment. The ultimate parameter of interest is the sign of $(\alpha_{HH} - \alpha_{HL}) - (\alpha_{HL} - \alpha_{LL})$. If negative, it provides evidence in favor of DARA and complementarity.

The specification for this test can be written as:

¹⁵These prior works often focused on testing complementarities in organization design as it pertains to firm performance. My method has been referred to in this literature as the “direct approach” or the “production approach” to testing complementarity. The direct approach does face endogeneity issues in that literature, because firm design and innovation strategy are often endogenous. However, endogeneity is not an issue in my analysis because my dimensions of interest have been randomized.

$$\begin{aligned}
\text{Adoption Choice}_{i,t} = & \beta_0 + \beta_1 \cdot \mathbb{1}(t = \text{High Context Uncertainty}) \cdot \mathbb{1}(t = \text{High Sampling Error}) \\
& + \beta_2 \cdot \mathbb{1}(t = \text{High Context Uncertainty}) \cdot \mathbb{1}(t = \text{Low Sampling Error}) \\
& + \beta_3 \cdot \mathbb{1}(t = \text{Low Context Uncertainty}) \cdot \mathbb{1}(t = \text{High Sampling Error}) \\
& + \beta_4 \cdot \mathbb{1}(t = \text{Low Context Uncertainty}) \cdot \mathbb{1}(t = \text{Low Sampling Error}) \\
& + u_i + \epsilon_{i,t}.
\end{aligned}$$

However, this formulation can be difficult to interpret. It features multiple events, each with a different coefficient whose sign and magnitude both matter for the final test. Instead, I formulate a variable combining multiple dummy terms from the above specification:

$$\begin{aligned}
\text{Difference}_{i,t} = & [\mathbb{1}(t = \text{High Context Uncertainty}) \cdot \mathbb{1}(t = \text{High Sampling Error}) \\
& - \mathbb{1}(t = \text{High Context Uncertainty}) \cdot \mathbb{1}(t = \text{Low Sampling Error})] \\
& - [\mathbb{1}(t = \text{Low Context Uncertainty}) \cdot \mathbb{1}(t = \text{High Sampling Error}) \\
& - \mathbb{1}(t = \text{Low Context Uncertainty}) \cdot \mathbb{1}(t = \text{Low Sampling Error})].
\end{aligned}$$

Using this new term, we can rewrite our specification as:

$$\text{Adoption Choice}_{i,t} = \beta_0 + \beta_1 \cdot \text{Difference}_{i,t} + u_i + \epsilon_{i,t}.$$

As a consequence of this specification, complementarity is now succinctly tested by whether $\beta_1 < 0$. The result of this estimation is in Table 3. While the coefficient is suggestive of complementarity, the standard error implies a low level of statistical significance. Similarly, a one tailed test yields a p-value of 0.177.

It is possible that these two sources of uncertainty are not complements. If so, this has important implications for those designing information provision: a reduction in context uncertainty will not increase the marginal value of reducing sampling error. However, two separate issues with the experiment suggest that the test is simply underpowered.

Because this specification involves testing an interaction, a much larger sample size is needed

Table 3: Regression output for test of complementarity

	(1)
Difference	-0.0150 (0.0162)
R ²	0.33917
Observations	6,336
Individual fixed effects	✓

for the same level of statistical power (i.e. [Gelman, Hill, and Vehtari 2020](#); [Gelman 2018](#)). One issue impacting power in this design is the desired effect size being measured. While context uncertainty could be substantially varied across rounds, switching from 25% to 75% of characters being of unknown types, the sampling error could not be as dramatically altered. This issue arose due to the combination of having a limited number of possible signal values, the limited number of tiles that fit on a tablet screen, and the requirement that the empirical distribution remain the same across types and context uncertainty rounds. Even if participants’ utility functions exhibit the curvature needed for DARA, they may still be locally linear given the small difference in sampling error.

Another issue impacting power is noise. The experiment relied heavily on enumerators providing proper instructions to participants, ensuring that all relevant details were provided. Poor instructions lead to a form of measurement error that also reduces statistical power. Out of twenty five enumerators, three were eventually dismissed due to their performance in providing instructions in the field. Table [A2](#) shows the impact of dropping data from these enumerators. This table provides an estimate that is both more negative and more statistically significant, with a one-tailed p-value of 0.062.

Overall, these results are suggestive of complementarity, but highlight the need for future work replicating this test. The ramifications for policymakers are substantial: understanding whether context uncertainty impacts how sampling error is perceived is important for designing information provision, especially when the cost of additional testing is high.

6.1 Heterogeneity Analysis

Thus far, analysis focused on determining whether context uncertainty impacts adoption among the entire sample of participant farmers. But how do farmer characteristics impact the relationship between context uncertainty and adoption?

Thus far, analysis focused on determining whether context uncertainty impacts adoption among the entire sample of participant farmers. But how do farmer characteristics impact the relationship between context uncertainty and adoption?

I examine heterogeneity by estimating a set of univariate regressions. Each regression explores how the difference in adoption level between high and low context uncertainty rounds varies with a different farmer demographic trait. The results are in Table 4.

These findings suggest little heterogeneity based on these traits. One exception is that farmers who attained education at least to the 6th standard or higher are more context averse. The latter has two potential interpretations. The first interpretation has to do with experiment comprehension. It could be that farmers with higher levels of education better understood the game’s design, and education therefore has little impact on context uncertainty in actual decision making. However, it is also possible that education helps farmers understand the risks of context uncertainty in all settings. Understanding this impact is left to future work.

Table 4: Context Uncertainty Aversion Heterogeneity Along Demographics

Characteristic					
Log Monthly Income	-0.03 (0.070)				
Age		(0.006) (0.006)			
Number of Dependents			0.00 (0.044)		
Number of Crops				0.00 (0.013)	
Married					0.24 (0.184)
Work Locally During Lean Season					0.10 (0.093)
Migration During Lean Season					-0.04 (0.118)
Early Adopter of Technology					0.13 (0.091)
Highest Education Attained					
Never Attended					—
Only pre-school					-0.34 (0.323)
Up to 5th standard					-0.16 (0.309)
Up to 8th standard					-0.52 (0.233)
Up to 10th standard					-0.53 (0.231)
Up to 12th standard					-0.67 (0.209)
Other					-0.61 (0.216)

Note: Standard errors in parenthesis. Context uncertainty data used from the high signal error module. Monthly income in INR and winsorized at 99% level. 12 observations dropped from marital regression for participants whose response was neither single nor married.

7 External Validity: Survey Evidence

Lab-in-the-field evidence from Section 6 documents that context uncertainty impacts behavior in a controlled setting. However, it is important to consider whether this behavior is consistent with farmers' behavior in practice. To evaluate this, the 1,571 farmers in my sample also participated in a survey. This section explores the responses from the survey, which included questions ranging from farmers' past experiences making technology adoption decisions, their perceptions about the informativeness of different sources (e.g. peers, extension agents), and the value they'd place on various hypothetical alternative extension programs.

Farmers' Perceptions of Prior Agricultural Advice

First, I asked each farmer about their past experience seeking agricultural advice. When asked whose advice they sought when making decisions about new agricultural technologies, 93% of farmers cited peers as an important source of information. These results, summarized in Figure 14, highlight that peers are the most important source of information for these farmers. An interesting observation within this data is that peers are also a more popular source of information than the local store, who also has a role as an information aggregator amongst consumers. This is potentially compatible with context uncertainty: when a farmer obtains information from peers directly, he knows the context behind each datum. By contrast, the information from the local store is aggregated. While certain dimensions of heterogeneity, such as weather, may have low variance amongst the store's customers, other dimensions, such as complementary inputs used on a plot, may have higher variance. Consequently, a farmer may have a harder time interpreting the store's aggregated signal.

Farmers were also asked to think of previous recommendations they received from their extension agent, and state how much it influenced their final decision. The results are reported in Figure 15. Over 60% of farmers reported at least some skepticism over the extension agent's recommendation. Those reporting any level skepticism were subsequently asked to list all applicable reasons for their concerns. Figure 16 groups their responses into two categories. The first category is reasons having to do with heterogeneous information, such as believing that your soil is different than the test plots or that the test plot operators had better training to take advantage of the technology. These concerns all pertain to technologies having heterogeneous returns, and believing that the positive signal shared by the extension agent won't necessarily translate into a farmer's own context. Of the 950




Figure 14: Cited Sources of Advice for New Agricultural Technologies

farmers who were asked the question, 590 of them (62%) stated at least one reason pertaining to information heterogeneity and 12% selected more than one reason in the category. These responses confirm that farmers within the sample have concerns about context uncertainty when evaluating information from their extension agent in real life.

What is the connection between citing these concerns within the survey and behavior within the experiment? Ideally, behavior would be a correlation between the two: farmers citing that context uncertainty reduced their trust in their extension agent would be those who show greater context uncertainty aversion in the experiment. Table 6 illustrates the results of a series of regressions, examining the correlation between the level of context uncertainty aversion in the experiment and each of the four context uncertainty reason cited within the survey.¹⁶ Concerns about different soil are significantly correlated, while the other three dimensions of context uncertainty are insignificant. Given the game design's focus on soil

¹⁶Individual fixed effects cannot be included in this regression, unlike the analysis of game performance in Section 6. The regression features differences in decisions across maps as a dependent variable and survey response to a single question as a independent variable. Both variables only have one observation per participant. In lieu of individual fixed effects, these regressions include behavioral controls for two pertinent randomizations within the experimental design. First, I control for whether, in the game module, the high context uncertainty map appeared before the low context uncertainty map for the participant. This dummy controls for anchoring issues happening within the module. Second, I control for whether the module used to measure context uncertainty for the regression appears to the participant as the fifth versus sixth module in the game.




Figure 15: Prior Extension Agent Influence




Figure 16: Cited Reasons for Skepticism of Prior Extension Recommendations

Table 6: How Context Uncertainty Aversion Correlates with Cited Concern for Heterogeneity

Characteristic				
Cite Soil Heterogeneity	-0.52 (0.199)			
Cite Heterogeneity of Other Inputs		0.23 (0.142)		
Cite Heterogeneity of Training			-0.17 (0.120)	
Cite Heterogeneity of Tech QC				-0.27 (0.249)
Module is Penultimate	0.25 (0.116)	0.26 (0.116)	0.26 (0.116)	0.26 (0.116)
High Context Uncertainty Map is First	-0.01 (0.116)	-0.01 (0.116)	-0.01 (0.116)	0.00 (0.116)

Note: Standard errors in parenthesis. Context uncertainty data used from the high signal error module. All regressions feature 950 observations. Individuals were dropped relative to full sample if they stated they completely trusted their extension agent’s past recommendations.

differences, by use of orange versus blue plot types, this is consistent with the fact that the soil was the most salient form of context uncertainty in the experiment.

Farmer’s Stated Preference for Test Plots

An implication of context uncertainty aversion is that farmers should be provided with greater information about test plot conditions. One rudimentary implementation involves providing soil test results of both a farmer’s own plot and the test plot, allowing farmers to compare their soil composition with those of a test plot. Compared to alternative methods to reduce context uncertainty, such as distributed test plots throughout a country, this approach is cost advantageous: it does not require a change to current infrastructure. In practice, this approach has two limitations, each of which can be incorporated into a simple extension of the model presented in this paper.

First, individuals may not understand how differences in an input dimension map to differences in a technology’s efficacy. While this mapping may be obvious in some circumstances (e.g. high potassium fertilizer is more beneficial for potassium deficient soil), it is less obvious in others (e.g. how much nitrogen is ideal for a new high yield variety corn seed). This is

not only true for the case of soil test results, but any for any input and technology mapping. An open question is how salient this concern is for individuals, such as the farmers in this sample.

Second, individuals may be concerned about unobserved dimensions of heterogeneity. For example, a technology’s efficacy may depend not only on soil content, but also on weather conditions, labor access, and a number of other inputs. The salience of these other dimensions was illustrated in Figure 16. Accordingly, these unobserved dimensions of heterogeneity may limit the efficacy of pairing centralized testing with detailed information about soil content.

To understand how salient these issues are to farmers, survey participants were asked to consider how influential recommendations based on four different hypothetical test plots would be on their future decisions, relative to the information they currently receive from their extension agent. Each example involved receiving data from a single test plot. The four cases varied along two dimensions.¹⁷ First, the test plot either had ‘very similar’ soil content or ‘substantially different’ soil content. Second, the test plot was either within 2 kilometers of their own plot (‘nearby’) or 100 kilometers away (‘far away’). Farmers’ preferences for these different cases is reported in Figure 17.




Figure 17: Farmer’s preferred test plot characteristic

These responses suggest that uncertainty about mapping and unobserved heterogeneity are

¹⁷To control for anchoring effects, respondents were randomized into four different groups. Each group corresponded to a distinct order for the four cases of possible test plots.

both salient to farmers. Consequently, there are potential large gains in influence from local testing. Notably, farmers’ desire for decentralized testing is not due to heterogeneous preferences. This puts forward an extension to Oates’ Decentralization Theorem (Oates 1972, 1993) whereby the public good, a test plot, should be supplied at the local level for reasons pertaining to heterogeneity and information, even when preferences are otherwise homogenous.

8 External Validity: Existing Literature

This section will be updated in the future to overview a variety of agricultural extension program designs in the literature and show that the history of which designs are effective is consistent with this paper’s model of context uncertainty.

9 Estimating Context Uncertainty Aversion

Throughout the paper, I have modeled context uncertainty as an additional form of risk. An implication of this modeling approach is that, because expected utility models agents as having a single risk aversion parameter, I have been unable to compare aversion to context uncertainty versus aversion to sampling error. However, my experimental design allows me to identify these parameters separately. This section will be expanded in the future to use the nested approach of Klibanoff, Marinacci, and Mukerji (2005) for modeling smooth ambiguity, developing a structural model that separately tests aversion to each form of uncertainty compare these parameters.

10 Conclusion

This paper presents evidence that context uncertainty is a significant factor in how individuals weigh signals when learning about new technologies. I show that this behavior exists in a well identified laboratory setting. I also provide evidence of external validity to farmers’ real world decisions about agricultural technologies and recommendations provided by extension agents.

One weakness of this paper’s approach is that it does not provide a new, actionable design for policymakers. An important direction for future field research is understanding how to

most efficiently design information provision to reduce context uncertainty. A companion project is currently under development to address this question.

Further, this project focuses on agricultural technology adoption in developing countries. However, examples of potential context uncertainty aversion can be found in a variety of settings, most recently in medicine (Alsan et al. 2022). However, there is evidence of agents exclusively caring about sampling error in other contexts, such as policymaking (Hjort et al. 2021). An open area for future research is understanding both which settings context uncertainty is important within and the reasons for these differences.

A related area of importance is better understanding the relationship between education and context uncertainty aversion. Do certain forms of education increase an individual’s ability to extrapolate information from sources with different characteristics? Do individuals highly educated in statistics, such as the policymakers in Hjort et al. (2021), learn to neglect context uncertainty?

The interaction between homophily and context uncertainty is another important area of exploration. Related questions have been explored theoretically, including Sethi and Yildiz (2016) and Baccara and Yariv (2013). However, several open questions remain about how the arrival of new technologies impacts homophily. For example, if an individual anticipates a high rate of new technologies arriving in the future, and expects to be uncertain about how to translate signals from other sources, are they more likely to form homophilous ties? Is this tendency reduced by improving education? Each of these questions poses interesting avenues for future theoretical and empirical research.

Overall, this project provides an important first step in a research agenda formally studying a common feature in previously effective information provision designs. Having been identified in a lab, there are now a number of useful extensions for future work. This project will hopefully also help information provision designers to better understand why certain designs have worked more effectively in the past and more efficiently modify future designs.

11 References

- Alismail, Sarah, and Hengwei Zhang. 2020. “Exploring and Understanding Participants’ Perceptions of Facial Emoji Likert Scales in Online Surveys: A Qualitative Study.” *ACM Transactions on Social Computing* 3 (2): 12:112:12. <https://doi.org/10.1145/3382505>.
- Alsan, Marcella, Maya Durvasula, Harsh Gupta, Joshua Schwartzstein, and Heidi L.

- Williams. 2022. “Representation and Extrapolation: Evidence from Clinical Trials,” October. <https://doi.org/10.3386/w30575>.
- Athey, Susan. 2002. “Monotone Comparative Statics Under Uncertainty*.” *The Quarterly Journal of Economics* 117 (1): 187–223. <https://doi.org/10.1162/003355302753399481>.
- Athey, Susan, and Scott Stern. 1998. “An Empirical Framework for Testing Theories about Complimentarity in Organizational Design,” June. <https://doi.org/10.3386/w6600>.
- Baccara, Mariagiovanna, and Leeat Yariv. 2013. “Homophily in Peer Groups.” *American Economic Journal: Microeconomics* 5 (3): 69–96. <https://doi.org/10.1257/mic.5.3.69>.
- Bar-Shira, Z., R.e. Just, and D. Zilberman. 1997. “Estimation of Farmers’ Risk Attitude: An Econometric Approach.” *Agricultural Economics* 17 (2-3): 211–22. <https://doi.org/10.1111/j.1574-0862.1997.tb00475.x>.
- Berardi, Michele. 2007. “Heterogeneity and Misspecifications in Learning.” *Journal of Economic Dynamics and Control* 31 (10): 3203–27. <https://doi.org/10.1016/j.jedc.2006.11.002>.
- Blackwell, David. 1951. “Comparison of Experiments.” *Proceedings of the Second Berkeley Symposium on Mathematical Statistics and Probability* 2 (January): 93–103. <https://projecteuclid.org/ebooks/berkeley-symposium-on-mathematical-statistics-and-probability/Proceedings-of-the-Second-Berkeley-Symposium-on-Mathematical-Statistics-and/chapter/Comparison-of-Experiments/bsmsp/1200500222>.
- . 1953. “Equivalent Comparisons of Experiments.” *The Annals of Mathematical Statistics* 24 (2): 265–72. <https://www.jstor.org/stable/2236332>.
- Cassiman, Bruno, and Reinhilde Veugelers. 2006. “In Search of Complementarity in Innovation Strategy: Internal r&d and External Knowledge Acquisition.” *Management Science* 52 (1): 68–82. <https://doi.org/10.1287/mnsc.1050.0470>.
- Clarke, Daniel J. 2016. “A Theory of Rational Demand for Index Insurance.” *American Economic Journal: Microeconomics* 8 (1): 283–306. <https://doi.org/10.1257/mic.20140103>.
- Dasaratha, Krishna, Benjamin Golub, and Nir Hak. 2021. “Learning from Neighbors about a Changing State.” <https://doi.org/10.48550/arXiv.1801.02042>.
- Dechezleprêtre, Antoine, Adrien Fabre, Tobias Kruse, Bluebery Planterose, Ana Sanchez Chico, and Stefanie Stantcheva. 2022. “Fighting Climate Change: International Attitudes Toward Climate Policies,” September. <https://doi.org/10.3386/w30265>.
- Dufo, Esther, and Emmanuel Saez. 2002. “Participation and Investment Decisions in a Retirement Plan: The Influence of Colleagues’ Choices.” *Journal of Public Economics* 85

- (1): 121–48. [https://doi.org/10.1016/S0047-2727\(01\)00098-6](https://doi.org/10.1016/S0047-2727(01)00098-6).
- Enke, Benjamin, and Thomas Graeber. 2019. “Cognitive Uncertainty.” <https://www.nber.org/papers/w26518>.
- Foster, Andrew D., and Mark R. Rosenzweig. 2010. “Microeconomics of Technology Adoption.” *Annual Review of Economics* 2 (September): 10.1146/annurev.economics.102308.124433. <https://doi.org/10.1146/annurev.economics.102308.124433>.
- Gelman, Andrew. 2018. “You Need 16 Times the Sample Size to Estimate an Interaction Than to Estimate a Main Effect.” <https://statmodeling.stat.columbia.edu/2018/03/15/need-16-times-sample-size-estimate-interaction-estimate-main-effect/>.
- Gelman, Andrew, Jennifer Hill, and Aki Vehtari. 2020. *Regression and Other Stories*. 1st edition. Cambridge New York, NY Port Melbourne, VIC New Delhi Singapore: Cambridge University Press.
- Giné, Xavier, Robert Townsend, and James Vickery. 2008. “Patterns of Rainfall Insurance Participation in Rural India.” *The World Bank Economic Review* 22 (3): 539–66. <https://doi.org/10.1093/wber/lhn015>.
- Gollier, Christian. 2011. “Portfolio Choices and Asset Prices: The Comparative Statics of Ambiguity Aversion.” *The Review of Economic Studies* 78 (4): 1329–44. <https://doi.org/10.1093/restud/rdr013>.
- Gollin, Douglas, and Christopher Udry. 2021. “Heterogeneity, Measurement Error, and Misallocation: Evidence from African Agriculture.” *Journal of Political Economy* 129 (1): 1–80. <https://doi.org/10.1086/711369>.
- Greis, Miriam, Hyunyoung Kim, Andreas Korge, Albrecht Schmidt, and Céline Coutrix. 2019. “SplitSlider: A Tangible Interface to Input Uncertainty.” In, edited by David Lamas, Fernando Loizides, Lennart Nacke, Helen Petrie, Marco Winckler, and Panayiotis Zaphiris, 493–510. *Lecture Notes in Computer Science*. Cham: Springer International Publishing. https://doi.org/10.1007/978-3-030-29390-1_27.
- Greis, Miriam, Hendrik Schuff, Marius Kleiner, Niels Henze, and Albrecht Schmidt. 2017. “Input Controls for Entering Uncertain Data: Probability Distribution Sliders.” *Proceedings of the ACM on Human-Computer Interaction* 1 (EICS): 3:13:17. <https://doi.org/10.1145/3095805>.
- Haaland, Ingar, Christopher Roth, and Johannes Wohlfart. Forthcoming. “Designing Information Provision Experiments.” *Journal of Economic Literature*, Forthcoming. <https://doi.org/10.1257/jel.20211658>.

- Haltiwanger, John, and Michael Waldman. 1985. "Rational Expectations and the Limits of Rationality: An Analysis of Heterogeneity." *The American Economic Review* 75 (3): 326–40. <https://www.jstor.org/stable/1814803>.
- Hjort, Jonas, Diana Moreira, Gautam Rao, and Juan Francisco Santini. 2021. "How Research Affects Policy: Experimental Evidence from 2,150 Brazilian Municipalities." *American Economic Review* 111 (5): 1442–80. <https://doi.org/10.1257/aer.20190830>.
- Karlan, Dean, Robert Osei, Isaac Osei-Akoto, and Christopher Udry. 2014. "Agricultural Decisions After Relaxing Credit and Risk Constraints *." *The Quarterly Journal of Economics* 129 (2): 597–652. <https://doi.org/10.1093/qje/qju002>.
- Kim, Hyeok, Dominik Moritz, and Jessica Hullman. 2021. "Design Patterns and Trade-Offs in Responsive Visualization for Communication." *Computer Graphics Forum* 40 (3): 459–70. <https://doi.org/10.1111/cgf.14321>.
- Klibanoff, Peter, Massimo Marinacci, and Sujoy Mukerji. 2005. "A Smooth Model of Decision Making Under Ambiguity." *Econometrica* 73 (6): 1849–92. <https://www.jstor.org/stable/3598753>.
- Koval, Morgane, and Yvonne Jansen. 2022. "Do You See What You Mean? Using Predictive Visualizations to Reduce Optimism in Duration Estimates." In, 119. CHI '22. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/3491102.3502010>.
- Krishnan, Pramila, and Manasa Patnam. 2014. "Neighbors and Extension Agents in Ethiopia: Who Matters More for Technology Adoption?" *American Journal of Agricultural Economics* 96 (1): 308–27. <https://doi.org/10.1093/ajae/aat017>.
- Liang, Annie. 2020. "Games of Incomplete Information Played by Statisticians," July. <https://doi.org/10.48550/arXiv.1910.07018>.
- Lu, Xuan, Wei Ai, Xuanzhe Liu, Qian Li, Ning Wang, Gang Huang, and Qiaozhu Mei. 2016. "Learning from the Ubiquitous Language: An Empirical Analysis of Emoji Usage of Smartphone Users." In, 770780. UbiComp '16. New York, NY, USA: Association for Computing Machinery. <https://doi.org/10.1145/2971648.2971724>.
- Magruder, Jeremy R. 2018. "An Assessment of Experimental Evidence on Agricultural Technology Adoption in Developing Countries." *Annual Review of Resource Economics* 10: 299–316.
- Mailath, George J., and Larry Samuelson. 2020. "Learning Under Diverse World Views: Model-Based Inference." *American Economic Review* 110 (5): 1464–1501. <https://doi.org/10.1257/aer.20190080>.

- Manski, Charles F. 2004. "Measuring Expectations." *Econometrica* 72 (5): 1329–76. <https://doi.org/10.1111/j.1468-0262.2004.00537.x>.
- McKenzie, David. 2022. "Why I Am Now More Cautious about Using or Recommending Matched Pair Randomization and Like Matched Quadruplets Instead." <https://blog.s.worldbank.org/impactevaluations/why-i-am-now-more-cautious-about-using-or-recommending-matched-pair-randomization>.
- Mobarak, Ahmed Mushfiq, and Mark R. Rosenzweig. 2012. "Selling Formal Insurance to the Informally Insured."
- Montiel Olea, José Luis, Pietro Ortoleva, Mallesh M Pai, and Andrea Prat. 2022. "Competing Models*." *The Quarterly Journal of Economics* 137 (4): 2419–57. <https://doi.org/10.1093/qje/qjac015>.
- Munshi, Kaivan. 2004. "Social Learning in a Heterogeneous Population: Technology Diffusion in the Indian Green Revolution." *Journal of Development Economics* 73 (1): 185–213. <https://doi.org/10.1016/j.jdeveco.2003.03.003>.
- Murphy, Kevin P. 2012. *Machine Learning: A Probabilistic Perspective*. Illustrated edition. Cambridge, MA: The MIT Press.
- Oates, Wallace E. 1972. *Fiscal Federalism*. Harcourt Brace Jovanovich.
- . 1993. "Fiscal Decentralization and Economic Development." *National Tax Journal* 46 (2): 237–43. <https://www.jstor.org/stable/41789013>.
- Qiao, Shan, Daniela B. Friedman, Cheuk Chi Tam, Chengbo Zeng, and Xiaoming Li. 2020. "Vaccine Acceptance Among College Students in South Carolina: Do Information Sources and Trust in Information Make a Difference?" <https://doi.org/10.1101/2020.12.02.20242982>.
- Rogers, Everett M. 2003. *Diffusion of Innovations, 5th Edition*. 5th edition. New York: Free Press.
- Rossi, Peter E., Greg M. Allenby, and Rob McCulloch. 2012. *Bayesian Statistics and Marketing*. John Wiley & Sons.
- Sajesh, V. K., and A. Suresh. 2016. "Public-Sector Agricultural Extension in India: A Note." *Review of Agrarian Studies* 6 (1): 116–31. <https://www.cabdirect.org/cabdirect/abstract/20163230852>.
- Salant, Yuval, and Josh Cherry. 2020. "Statistical Inference in Games." *Econometrica* 88 (4): 1725–52. <https://doi.org/10.3982/ECTA17105>.
- Sethi, Rajiv, and Muhamet Yildiz. 2016. "Communication With Unknown Perspectives." *Econometrica* 84 (6): 2029–69. <https://doi.org/10.3982/ECTA13320>.

- Shannon, Chris. 1995. “Weak and Strong Monotone Comparative Statics.” *Economic Theory* 5 (2): 209–27. <https://www.jstor.org/stable/25054834>.
- Suri, Tavneet. 2011. “Selection and Comparative Advantage in Technology Adoption.” *Econometrica* 79 (1): 159–209. <https://doi.org/10.3982/ECTA7749>.
- Suri, Tavneet, and Christopher Udry. 2022. “Agricultural Technology in Africa.” *Journal of Economic Perspectives* 36 (1): 3356.
- Takahashi, Kazushi, Yukichi Mano, and Keijiro Otsuka. 2019. “Learning from Experts and Peer Farmers about Rice Production: Experimental Evidence from Cote d’Ivoire.” *World Development* 122 (October): 157–69. <https://doi.org/10.1016/j.worlddev.2019.05.004>.

A The Impact of Priors on Measuring Adoption

B Design of Behavioral Modules

C Results from Behavioral Modules

D Randomization Algorithm

E Tables for Primary Hypotheses

Table A1: Regression output for effect of increasing context uncertainty on adoption intensity

	Low Sampling Error (1)	High Sampling Error (2)	Pooled (3)
High Context Uncertainty	-0.2929*** (0.0461)	-0.3529*** (0.0456)	-0.3229*** (0.0325)
R ²	0.81466	0.82333	0.34465
Observations	3,168	3,168	6,336
Individual fixed effects	✓	✓	✓

Table A2: Regression output for test of complementarity without dropped enumerators

	(1)
Difference	-0.0253 (0.0164)
R ²	0.34177
Observations	5,936
Individual fixed effects	✓

F Proofs

Proof of Proposition 3.1

Proof. The distributions of each signal is centered at $\theta + \gamma_i$. Assume without loss of generality that the extension agent's signal, S_E^A , has greater variance. We can parameterize the difference in variance between the signals as

$$\delta = (\sigma_E^2 + (\sigma_E^\gamma)^2) - (\sigma_j^2 + (\sigma_j^\gamma)^2).$$

Next, consider a signal $s_k^A = s_j^A + \zeta$ where $\zeta \sim \mathcal{N}(0, \delta)$ and independent of all other random variables composing s_j^A . Because ζ is independent Gaussian noise with mean 0, s_k^A is considered a *garbling* of s_j^A . This is equivalent to the agent preferring s_j^A over s_k^A by Blackwell's Theorem (Blackwell 1951, 1953). Next, observe that $s_k^A \sim_A \mathcal{N}(\theta + \gamma_i, \sigma_j^2 + (\sigma_j^\gamma)^2 + \delta)$. By construction of δ , this implies $s_k^A \sim_A \mathcal{N}(\theta + \gamma_i, \sigma_E^2 + (\sigma_E^\gamma)^2)$, and so s_k^A and s_E^A are equivalent signals about θ_A . By transitivity, the agent prefers s_j^A over s_E^A .

□

Proof of Proposition 4.1

Proof. Note that the decision problem in Equation 10 can be expanded to

$$\alpha^* = \arg \max_{\alpha} \int_{\theta_A} \int_{\epsilon_{A,1}} \int_{\epsilon_{A,0}} \int_{\nu_A} u(1 + \nu_A + \epsilon_{A,0} + \alpha(\theta_A + \epsilon_{A,1} - \epsilon_{A,0})) dF_{\nu_A} dF_{\epsilon_{A,0}} dF_{\epsilon_{A,1}} dF_{\theta_A}.$$

Integrating over the distributions of our independent exogenous shocks $\epsilon_{A,0}$, $\epsilon_{A,1}$, and ν_A yields

$$\alpha^* = \arg \max_{\alpha} \int_{\theta_A} E[u(1 + \nu_A + \epsilon_{A,0} + \alpha(\theta_A + \epsilon_{A,1} - \epsilon_{A,0})) | \epsilon_{A,0}, \epsilon_{A,1}, \nu_A] dF_{\theta_A}.$$

The precision of a univariate distribution is simply the inverse of its variance. To compress notation and focus on our parameters α and θ_A , we let

$$v(Y_A(\alpha, \theta_A)) \equiv E[u(1 + \nu_A + \epsilon_{A,0} + \alpha(\theta_A + \epsilon_{A,1} - \epsilon_{A,0})) | \epsilon_{A,0}, \epsilon_{A,1}, \nu_A]$$

and rewrite our maximization problem as

$$\alpha^* = \arg \max_{\alpha} \int_{\theta_A} v(Y_A(\alpha, \theta_A)) f_{\theta_A}(t; 1/\tilde{\sigma}_0^2) dt$$

where $f_{\theta_A}(t; 1/\tilde{\sigma}_0^2)$ denotes the probability density function of θ_A as a function of precision in the posterior belief, $1/\tilde{\sigma}_0^2$.

Our proof will take three steps. First, we will show that α^* is nondecreasing in the overall precision of the posterior $1/\tilde{\sigma}_0^2$. Because precision is the inverse of variance, this immediately implies that α^* is nonincreasing in variance $\tilde{\sigma}_0^2$. Finally, we will show that this implies it is also nonincreasing in context uncertainty $(\sigma_j^\gamma)^2$ from any signal.

Recall that u is nondecreasing and concave. The conditional expectation is a linear operator. Thus, these properties are preserved; v is also nondecreasing and concave. This observation will be used shortly.

Note that Y_A is a linear function of the random variable θ_A . Consequently, for two values x'' and x' , if the distribution $F_{\theta_A}(x')$ is a mean preserving spread of $F_{\theta_A}(x'')$, then $F_{Y_A}(\alpha, x')$

is a mean preserving spread of $F_{Y_A}(\alpha, x'')$. Because F_{θ_A} is parameterized by its precision, this is true whenever $x'' > x'$.

If a distribution B is a mean preserving spread of another distribution A , then A second-order stochastically dominates B . Further, a distribution A dominates B in the the second order if and only if $E[u(A)] \geq E[u(B)]$ for all functions u that are nondecreasing and concave.

We have now proven three facts. First, our function v is nondecreasing and concave. Second, for $x'' > x'$, $F_{Y_A}(\alpha, x')$ is a mean preserving spread of $F_{Y_A}(\alpha, x'')$. Third, if a distribution B is a mean preserving spread of another distribution A , then $E[v(A)] \geq E[v(B)]$ for all functions v that are non-decreasing and concave. Together, these imply that

$$E[v(F_{Y_A}(\alpha, x''))] \geq E[v(F_{Y_A}(\alpha, x'))].$$

Next, we borrow standard notation from expected utility theory and state that $U(\alpha, x) = E[v(F_{Y_A}(\alpha, x))]$. Therefore, we have proven that $U(\alpha, x'') - U(\alpha, x') \geq 0$ whenever $x'' > x'$. Shannon (1995) proves that α^* is nondecreasing in x if $U(\alpha, x)$ obeys the *weak single crossing property* in (α, x) : for any $\alpha'' > \alpha'$ and $x'' > x'$,

$$U(\alpha', x'') \geq U(\alpha', x') \implies U(\alpha'', x'') \geq U(\alpha'', x').$$

The weak single crossing property is trivially true for U in (α, x) , because we proved that $U(\alpha, x'') \geq U(\alpha, x')$ whenever $x'' > x'$, regardless of α . Thus, we have proven that the agent's optimal share of adoption α^* is nondecreasing in the overall precision of the posterior, $1/\tilde{\sigma}_0^2$.

Because the precision is the inverse of the variance, it immediately follows that the agent's optimal share of adoption α^* is nonincreasing in the overall variance of the posterior, $\tilde{\sigma}_0^2$. Finally, it is immediate from Equation 8 that, all else held constant, $\tilde{\sigma}_0^2$ is increasing in the context uncertainty $(\sigma_j^\gamma)^2$ regarding any individual j . Because this relationship is monotonic, our comparative static is preserved: the agent's optimal share of adoption α^* is nonincreasing in the context uncertainty $(\sigma_j^\gamma)^2$ regarding any individual j .

□

Proof of Proposition 4.2

Proof. To show that, for an agent with DARA preferences, context uncertainty $(\sigma_j^\gamma)^2$ and sampling uncertainty σ_j^2 are decision complementary, we must show that the optimal level of adoption α^* has the property

$$\alpha^*(\sigma_j^{\gamma''}, \sigma_j'') - \alpha^*(\sigma_j^{\gamma'}, \sigma_j') \geq \alpha^*(\sigma_j^{\gamma'}, \sigma_j'') - \alpha^*(\sigma_j^{\gamma'}, \sigma_j')$$

for all $\sigma_j^{\gamma''} \geq \sigma_j^{\gamma'}$ and $\sigma_j'' \geq \sigma_j'$.

If the agent's utility is thrice differentiable, then DARA implies $U''' > 0$. If so, the third partial derivative $\frac{\partial^3 U}{\partial \alpha \partial \sigma_j \partial \sigma_j^\gamma} > 0$. Our inequality immediately follows.

If the agent's utility is not thrice differentiable, we can still show decision complementarity. Assume for the sake of contradiction that our inequality did not hold. Then for some $\sigma_j^{\gamma''} \geq \sigma_j^{\gamma'}$ and $\sigma_j'' \geq \sigma_j'$, we have

$$\alpha^*(\sigma_j^{\gamma''}, \sigma_j'') - \alpha^*(\sigma_j^{\gamma'}, \sigma_j'') < \alpha^*(\sigma_j^{\gamma''}, \sigma_j') - \alpha^*(\sigma_j^{\gamma'}, \sigma_j').$$

But, consider the impact of context uncertainty on any given realization of signal noise. Because the agent's preferences exhibit DARA, the agent is more risk averse at lower realizations than at higher realizations. Consequently, at high realizations of signal noise, the increase in context uncertainty will impact the agent's optimal adoption choice less at lower realizations.

Next, we can interpret an increase in signal noise as a mean preserving spread moving weight away from the mean and towards the tails, increasing the likelihood of extreme values. Because of DARA, this means that the impact, measured as the difference in optimal adoption as context uncertainty changes, is greater below the mean than above the mean.

Above mean realizations of signal noise, the impact of context uncertainty at the margin is reduced by the mean preserving spread. This is because the weight of the distribution is going towards higher signal error realizations where risk aversion is lower and context uncertainty has less impact. Below the mean, weight is being moved towards lower signal error realizations where risk aversion is higher and context uncertainty has more impact.

However, the mean preserving spread is more impactful below the mean than above the mean because DARA implies there is more concavity below the mean. Thus, a mean preserving spread redistributes weight towards realizations where changes in context uncertainty have greater impact.

It follows that our inequality cannot hold.

□